

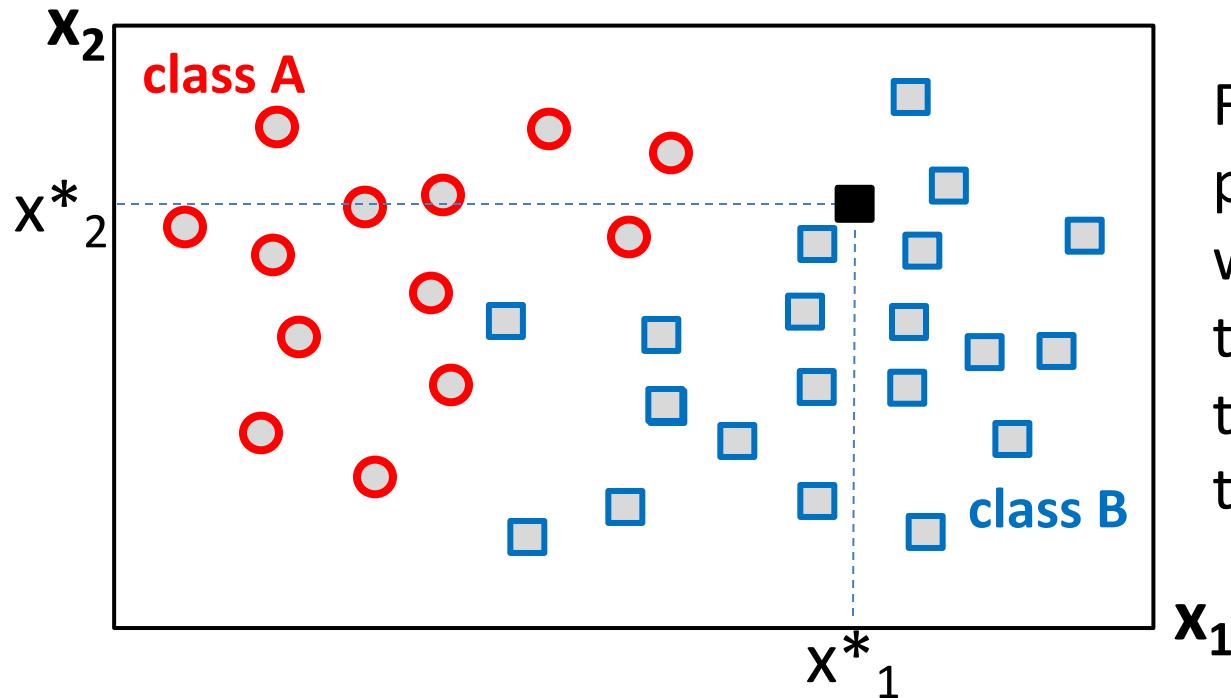
Machine Learning

Lecture 08 k Nearest Neighbors

Characteristics of kNN

- k-Nearest Neighbors (**kNN**) for predictive modeling
 - A simple, supervised, **instance-based** algorithm
 - **Non-linear** and **non-parametric**
 - A lazy learner (no training)
 - Very intuitive (high interpretability)
- Can handle:
 - Mixed data types (continuous and categorical, but careful there!)
 - Linearly inseparable data
 - Multi-class problems
- When to consider: Recommender systems, fault and intrusion detection, low-dimensional datasets, document categorization, target marketing, great for remote sensing data.

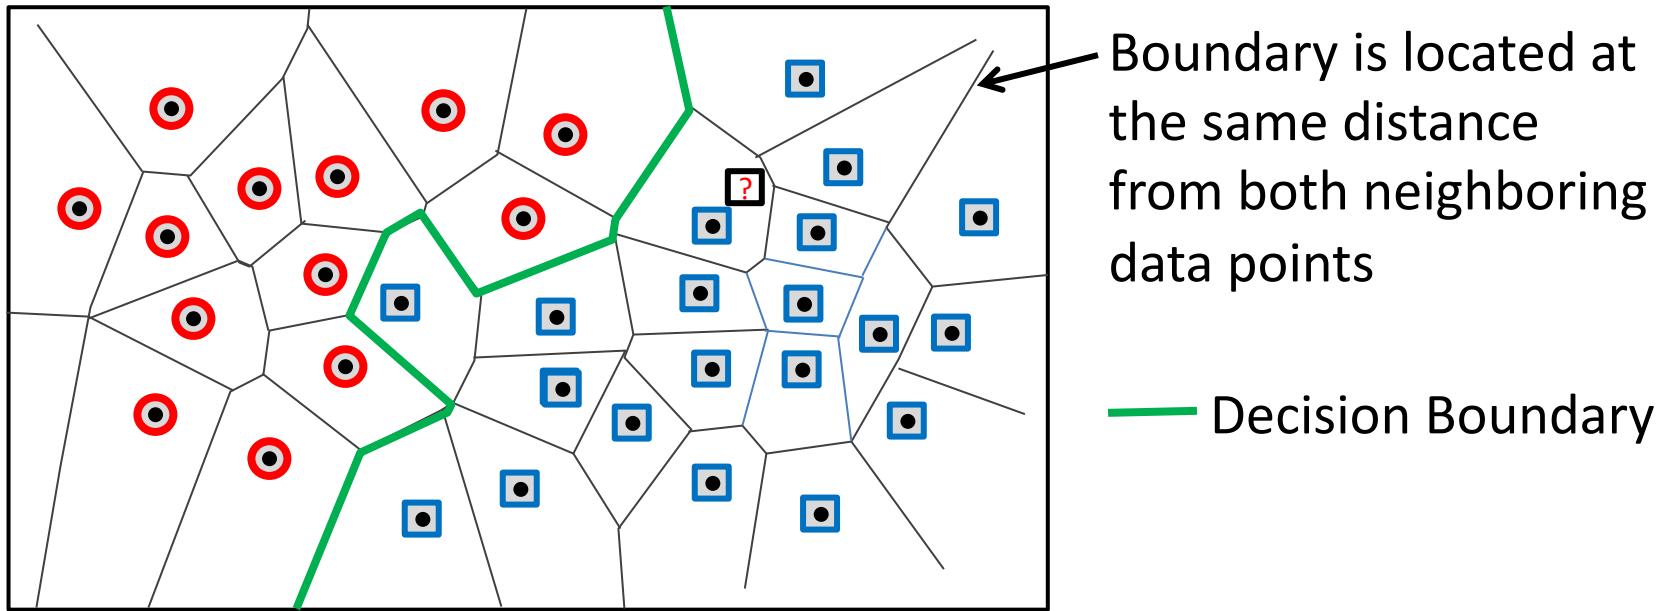
- Which class does ■ belong to?



For a given set of points (x_1^*, x_2^*) , which one of the two classes does this point belong to?

- Key point: **nearby points \Rightarrow same class**
- Based on the idea that closer two objects are in space, the more similar they are
- Need a good distance function to establish similarity

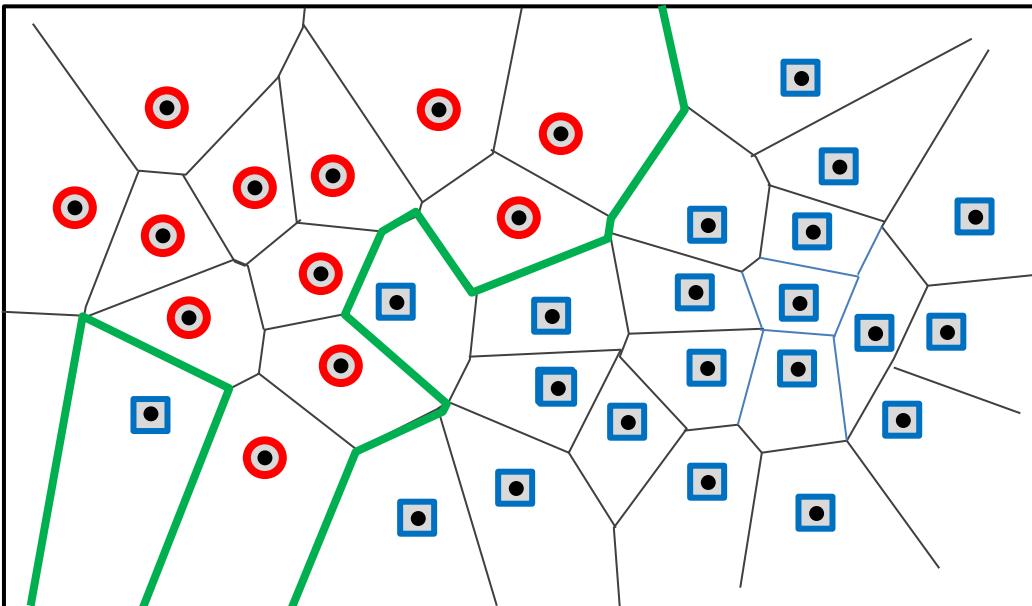
- Voronoi Diagram



- Partitions the whole space into cells and finds the most similar training data point to \square and predicts class
- A complicated decision boundary (nonlinear—composed of broken straight lines) which separates the classes well

Sensitivity to outliers

- What happens when you change the class label of a single data point?



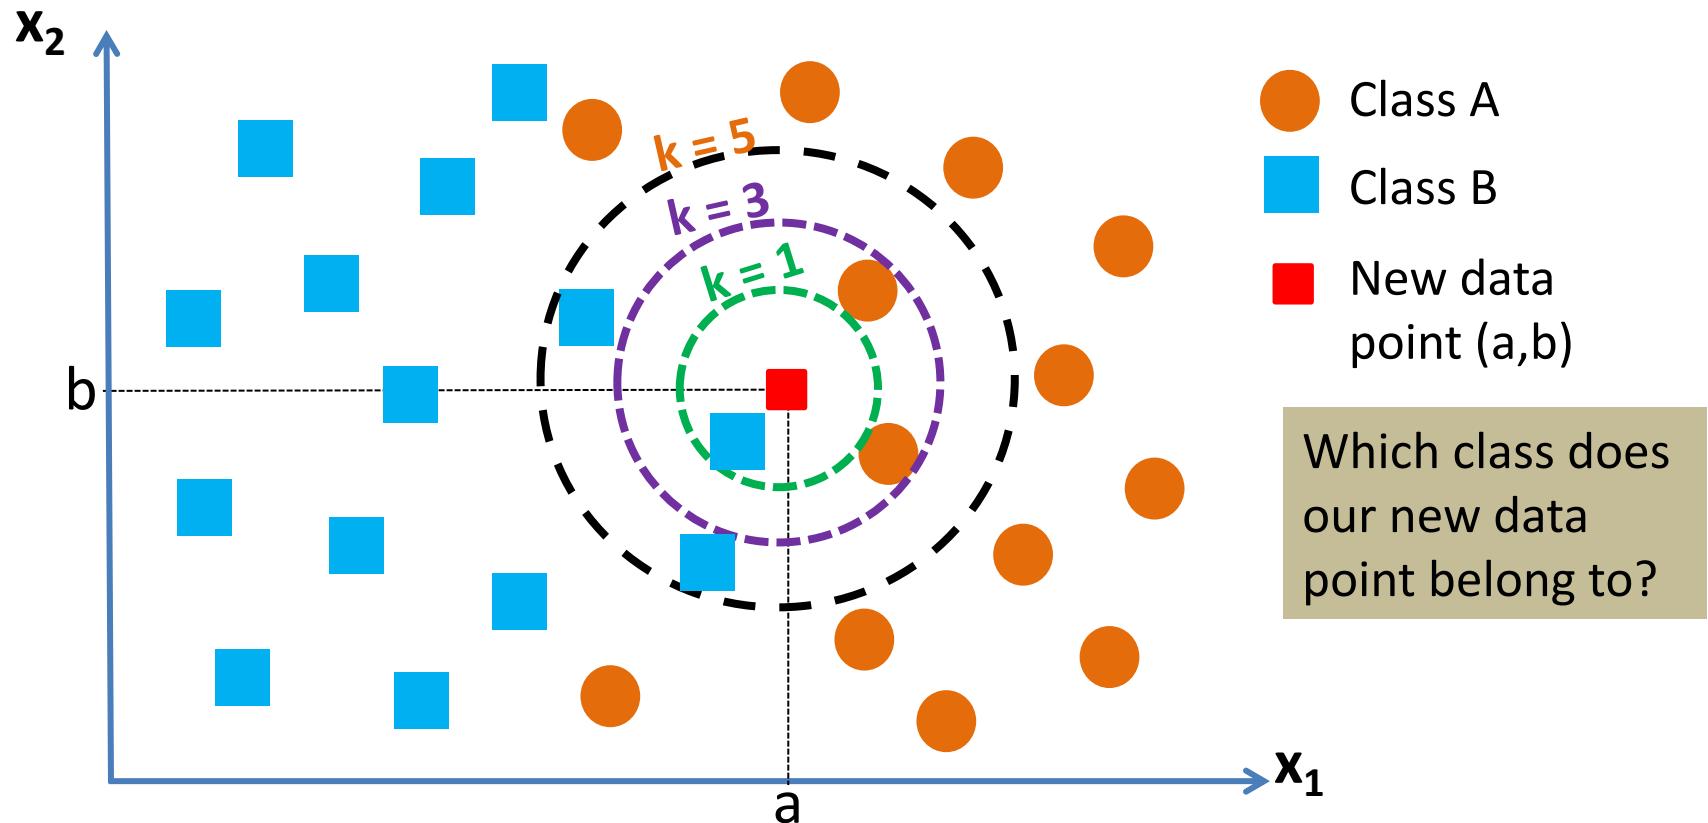
Creates a dramatic change in the decision boundary which is not desirable.

We don't want small changes in the data to have large effects on the prediction.

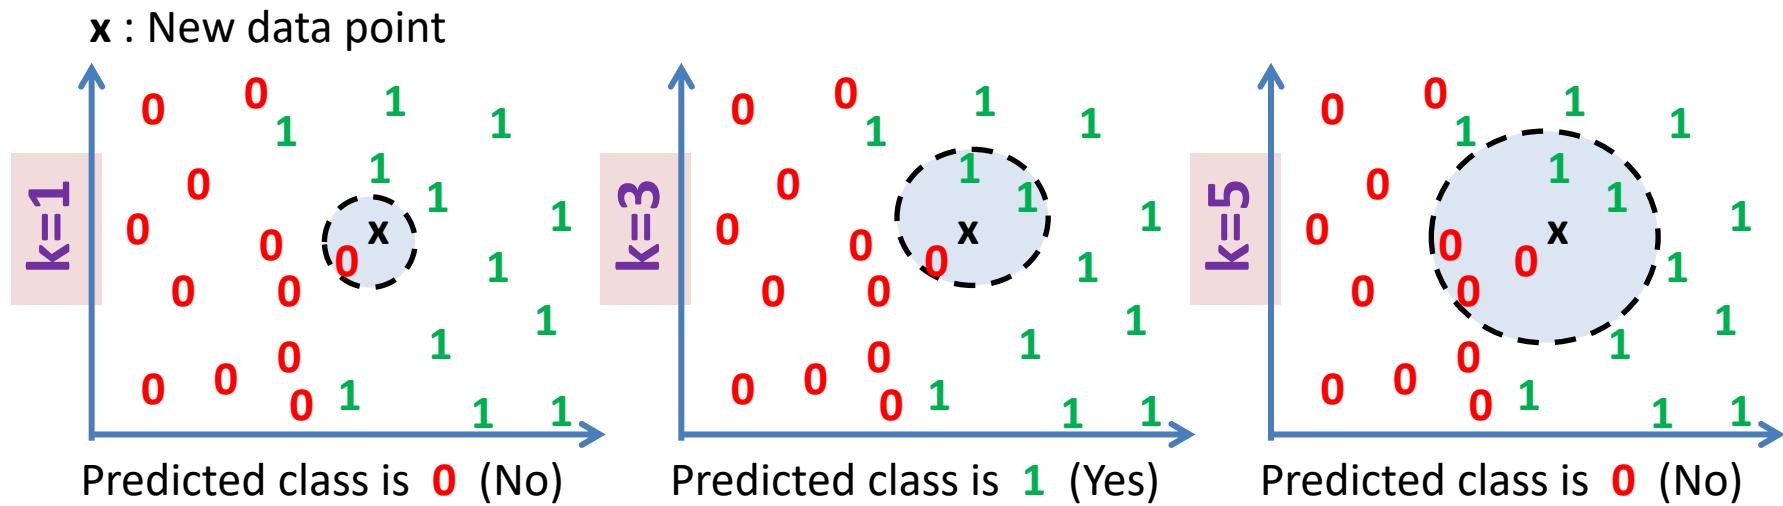
- Mislabeled training data have a compound effect on the class prediction. Why?
- No conditional probability in the form of $p(y|x)$, i.e., it's not sensitive to class prior.
- Is there a way to make kNN sensitive to data counts?

Nearest Neighbors algorithm

- How about using more than one nearest neighbor for class prediction?
- Spot the "k" most similar neighbors to the new data and count class labels
- Metric for quantifying similarity : "**distance**"



How many neighbors?



k is an odd value for no ties (if binary class)

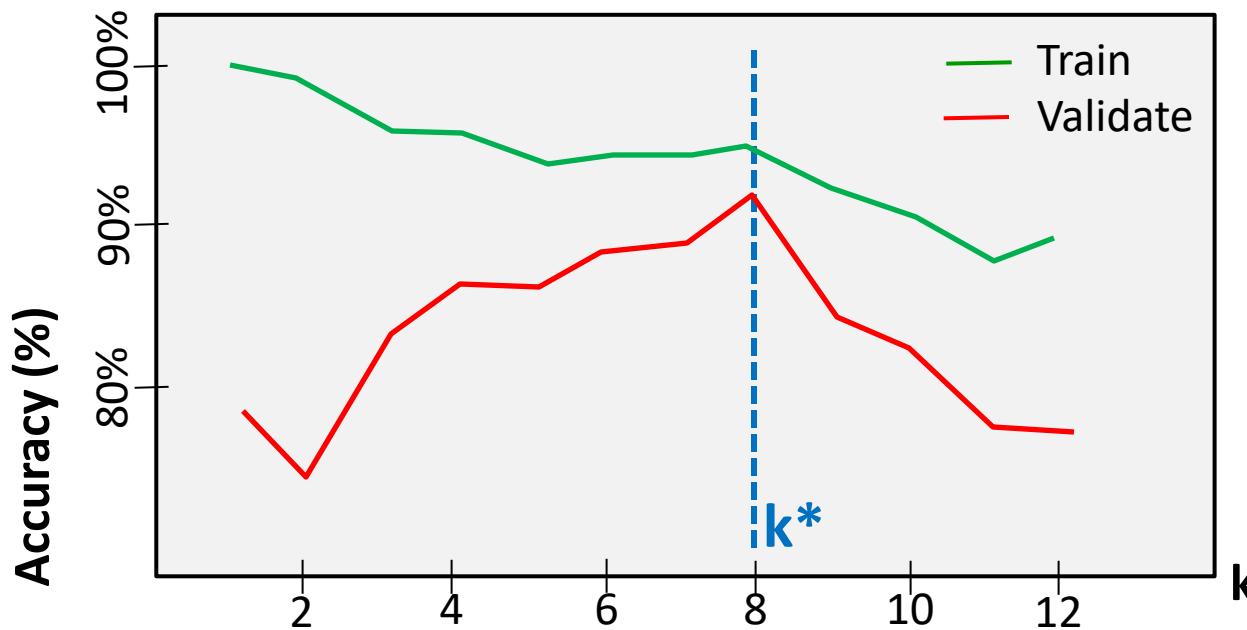
As **k** gets larger (larger neighborhood) → New data point gets classified as the most probable class (large bias, inaccurate estimation)

As **k** gets smaller (smaller neighborhood) → Sensitive to noise and highly variable (unreliable estimation as a consequence of overfitting)

A probabilistic estimate could be given based on the classes of the neighbors if you want to assign a score to the prediction (other than a simple Yes or No)

How many neighbors? – cont'd

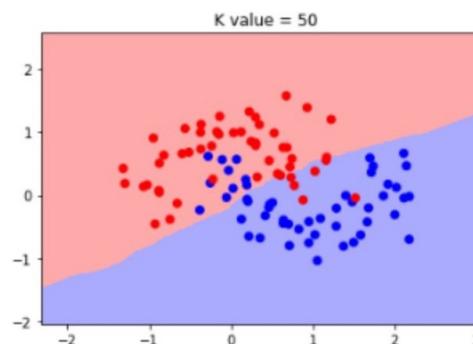
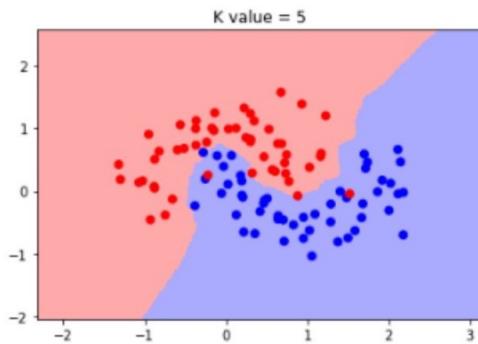
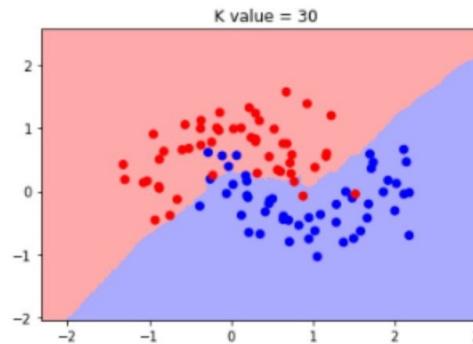
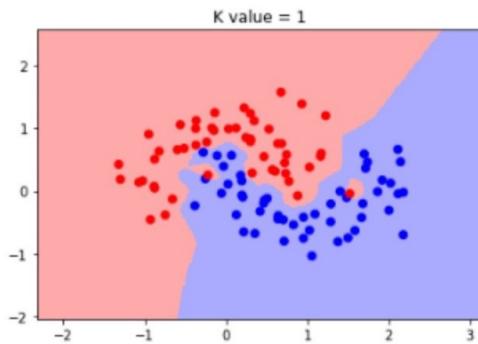
- **Methodology:**
 - Split the whole data into train and validation sets
 - Select a range of "k" for the number of neighbors
 - For each "k", do a cross validation and compute the error (or accuracy) for both train and validation sets
 - Plot the train-validate error over the range of k's



Pick k^* which gives the lowest validation error (for the best generalization performance)

How many neighbors? – cont'd

- How does k affect the performance of the model?
 - The bigger the k , the smoother the decision boundary. If the difference in the CV errors is negligible for increasing k , a larger value of k may be chosen if computational expense isn't an issue.
 - If the CV error doesn't start to rise again, it could mean that the features are not informative and a constant output is the best it can do.

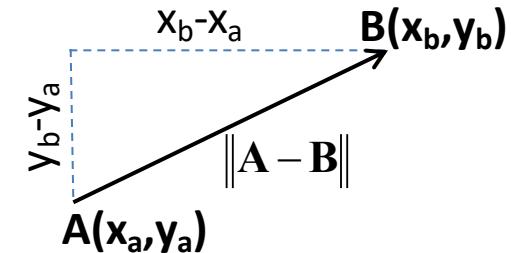


Large k will over-smooth ignoring local structure

Influence of neighbors on prediction

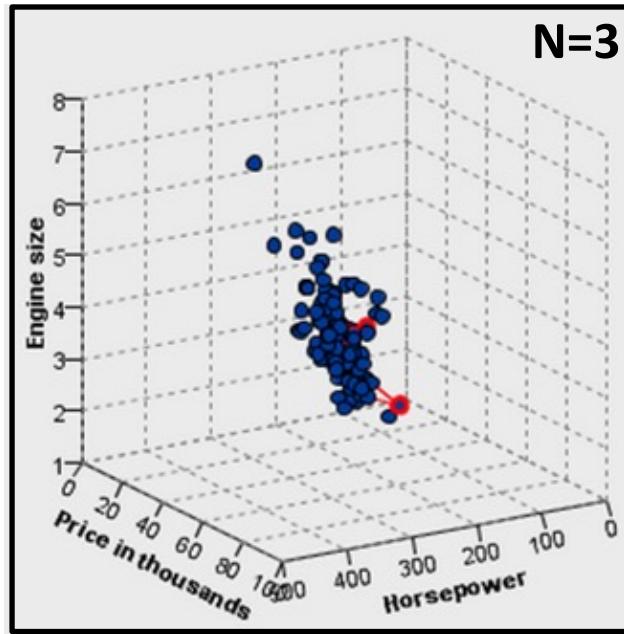
- Distance between neighbors?
- A key component of kNN algorithm
 - A metric for measuring how similar instances are
 - Has a strong impact on classification performance
- k neighbors that are closest to the new data point “ x ” are determined by a distance measure
- Measures of distances used in kNN:
- Euclidean distance
 - Euclidean (L2 norm) distance in 2D (valid for numerical variables)

$$d(\mathbf{A}, \mathbf{B}) = \sqrt{(x_b - x_a)^2 + (y_b - y_a)^2}$$



Influence of neighbors on prediction

Euclidean distance
in higher dimensions



Brand	Engine size	Horsepower	Price
Ford Torino	302	140	15000
AMC rebel	304	150	16500
Buick skylark	350	165	19000

$$d(\text{ford,amc}) = \sqrt{(302 - 304)^2 + (140 - 150)^2 + (15000 - 16500)^2} \\ = 1500$$

$$d(\text{ford,buick}) = \sqrt{(302 - 350)^2 + (140 - 165)^2 + (15000 - 19000)^2} \\ = 4000.4$$

Source: <http://www.datasciencecentral.com/profiles/blogs/introduction-to-the-k-nearest-neighbor-knn-algorithm>

- Euclidean distance in "N" dimensions:

$$d(\mathbf{A}, \mathbf{B}) = \sqrt{(d_{1,A} - d_{1,B})^2 + (d_{2,A} - d_{2,B})^2 + \dots + (d_{N,A} - d_{N,B})^2} \\ = \sqrt{\sum_{i=1}^N (d_{i,A} - d_{i,B})^2} \quad (\text{valid for continuous variables})$$

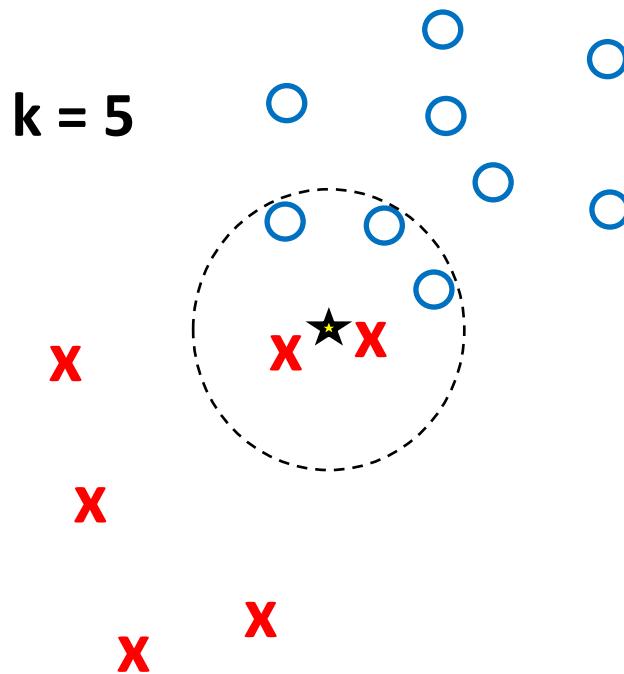
Influence of neighbors on prediction

- Once we find \mathbf{x} 's nearest neighbors using the distance function, neighbors vote to predict \mathbf{x} 's class:
- Method 1: **Majority voting**
 - All votes are equal and majority (democracy) wins
- Method 2: **Weighted voting**
 - Closer neighbors get higher votes (shareholder democracy)
 - Distance-based weight: A neighbor's vote is the inverse of its distance to \mathbf{x} :
$$vote(\mathbf{x}_i) = 1/d(\mathbf{x}_i, \mathbf{x})$$
 - Heuristic: Based on domain-specific characteristics of neighbors

Influence of neighbors on prediction

- Majority Voting vs Weighted Voting mechanisms

Majority-voting
(uniform in sklearn)

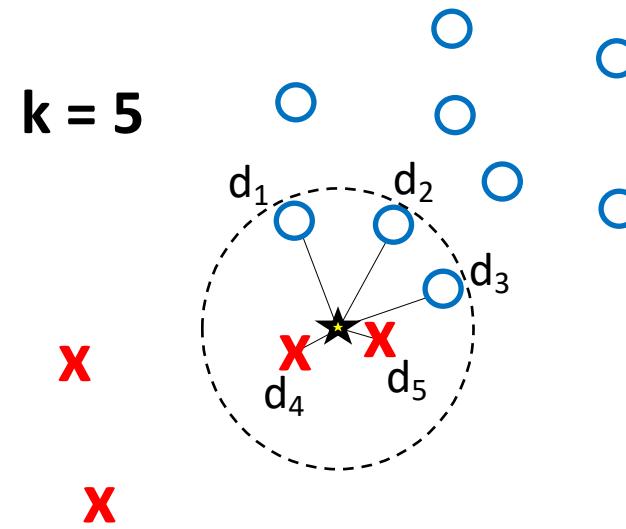


O : 3
X : 2



Prediction: O

Weighted-voting
(distance in sklearn)



$$\begin{aligned} \text{O} &: \frac{1}{d_1} + \frac{1}{d_2} + \frac{1}{d_3} \\ \text{X} &: \frac{1}{d_4} + \frac{1}{d_5} \end{aligned}$$



Prediction: X

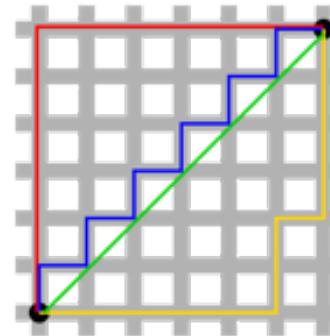
Other measures of distance

- **Minkowski distance (m-norm)**

$$\begin{aligned} d(\mathbf{A}, \mathbf{B}) &= \sqrt[p]{|d_{1,A} - d_{1,B}|^p + |d_{2,A} - d_{2,B}|^p + \dots + |d_{N,A} - d_{N,B}|^p} \\ &= \sqrt[p]{\sum_d |d_{i,A} - d_{i,B}|^p} = \left(\sum_d |d_{i,A} - d_{i,B}|^p \right)^{1/p} \end{aligned}$$

- If $p=0 \Rightarrow$ **Hamming distance** (later)
- If $p=1 \Rightarrow$ **Manhattan distance** (L1 norm)
 - Suitable for grid-like paths such as maps (usually preferred in the presence of many features – high dimensional data)

$$d(\mathbf{A}, \mathbf{B}) = \sum_d |d_{i,A} - d_{i,B}|$$

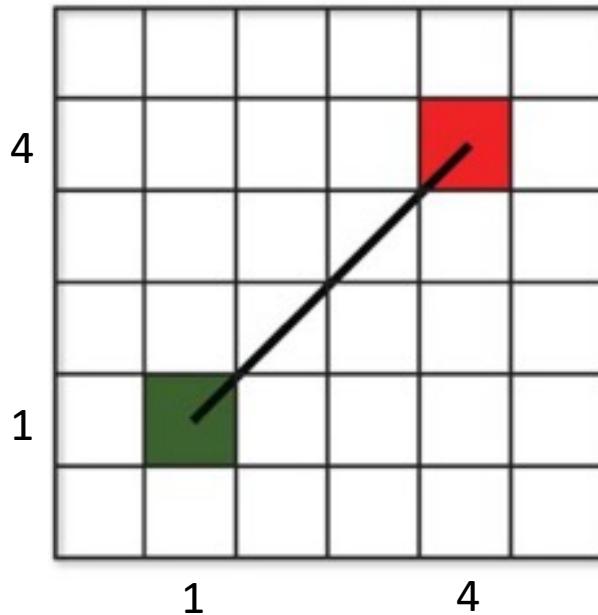


Taxicab or city block distance

- If $p=2 \Rightarrow$ **Euclidean distance** (L2 norm)

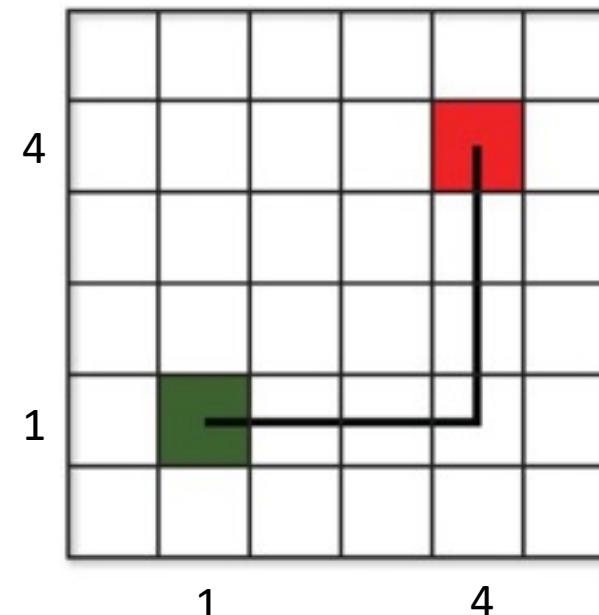
Other measures of distance

- Euclidean vs Minkowski



Euclidean distance between A and B:
Straight line distance (most
commonly used metric – default one)

$$d(\mathbf{A}, \mathbf{B}) = \sqrt{(4-1)^2 + (4-1)^2} = 4.24$$



Manhattan distance between A and B:
The distance between two points is
the sum of the absolute differences of
their Cartesian coordinates, i.e.,

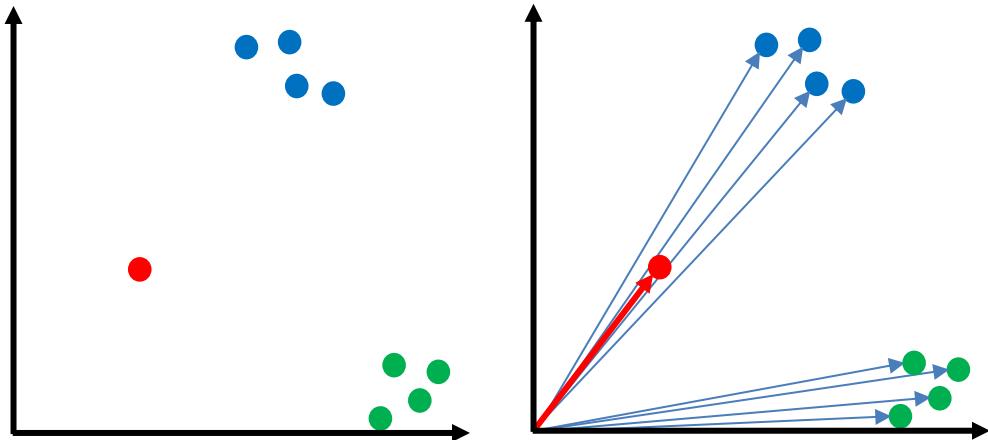
$$d(\mathbf{A}, \mathbf{B}) = \sum_{i=1}^2 |x_i - y_i| = |4-1| + |4-1| = 6$$

Other measures of distance – cont'd

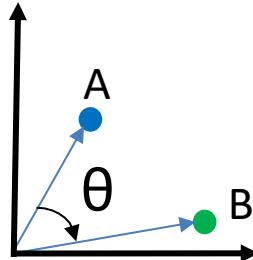
- **Distance based on Cosine Similarity**

- A class distribution as given below will not lead to a good separation between the classes using Euclidian distance.

Instead, we can use
a metric based on
Cosine similarity
(measuring the
angle between
the class vectors)



- So a Cosine similarity between two data points is given as:



$$\cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|} = \frac{\sum A_i B_i}{\sqrt{\sum A_i^2} \sqrt{\sum B_i^2}}$$

$$\text{Cos distance} = 1 - \cos(\theta)$$

The closer the class vectors are by angle, the higher is the Cosine Similarity (Cos theta), thus smaller the distance.

Other measures of distance – cont'd

- Cosine similarity is very useful when we are interested in the orientation but not the magnitude of the vectors.
- Frequently used in text analytics where the similarity between documents is investigated.

Example:		I	think	therefore	am	Can	you	don't	know	who
1	I think, therefore I am	2	1	1	1	0	0	0	0	0
2	Can you think?	0	1	0	0	1	1	0	0	0
3	I don't think, therefore I don't know who I am	3	1	1	1	0	0	2	1	1

$$d_{12} = 1 - \frac{1}{\sqrt{2^2 + 1 + 1 + 1} \sqrt{1 + 1 + 1}} = 0.782$$

$$d_{13} = 1 - \frac{6 + 1 + 1 + 1}{\sqrt{2^2 + 1 + 1 + 1} \sqrt{3^2 + 1 + 1 + 1 + 2^2 + 1 + 1}} = 0.198$$

$$d_{23} = 1 - \frac{1}{\sqrt{1 + 1 + 1} \sqrt{3^2 + 1 + 1 + 1 + 2^2 + 1 + 1}} = 0.864$$

Sentence 1 and 3
are closest.

Other measures of distance – cont'd

- What if the features are categorical?
- Categorical attributes
 - Univalent: Single level (an object could be black or white but not both)
 - Nominal variables
 - Ordinal variables
 - Multivalent: A specific object may contain multiple levels (a movie genre could be both Action and Comedy at the same time)
- There are several distance functions measuring dissimilarity for categorical attributes.

Distance for categorical attributes

- **Nominal attributes**

- "Hamming distance" for binary variables ($p=0$):
- Example: {male, female}, {green, blue, red}

$$d_H(\mathbf{x}, \mathbf{y}) = \sum_{k=1}^N |x_k - y_k| \cdot \begin{cases} x = y \Rightarrow D = 0 \\ x \neq y \Rightarrow D = 1 \end{cases}$$

Person1	Person2	Distance
Male	Female	1
Female	Female	0

- **Method I (Simple Matching)**: (A more general formula including multi-class for number of categories > 2):

$$d(i, j) = \frac{p - m}{p} = \frac{\# \text{ of mismatches}}{\text{total } \# \text{ of variables}}$$

m: Number of matches

$$d(\text{male}, \text{female})=?$$

$(1,0) \Leftrightarrow (0,1) \Rightarrow \text{zero matches} : m = 0, p = 2 \Rightarrow d(\text{male}, \text{female}) = 1$

- **Method II**: If many levels in an attribute, create a new binary attribute for each of the M nominal states

Distance for categorical attributes – cont'd

- **Ordinal attributes**

- An ordinal variable can be discrete or continuous
- Order (rank) is important: {small, medium, large}
- Can be treated like interval-scaled
 - Replace an ordinal variable by its rank $r_{ik} \in \{1, 2, \dots M_k\}$
 - Map the range of each variable onto [0,1] by replacing the i^{th} object in k^{th} variable by:

$$z_{ik} = \frac{r_{ik} - 1}{M_k - 1}$$

Data set = {**cold** , **cool** , **warm** , **hot** }
cold = $(1-1)/(4-1) = 0$ **cool** = $(2-1)/(4-1) = 1/3$
warm = $(3-1)/3 = 2/3$ **hot** = $(4-1)/3 = 1$
Example: {**cold**: 0, **cool**: 1/3, **warm**: 2/3, **hot**: 1}

- Distance between **warm** and **hot**: $1 - 2/3 = 1/3$
- Distance between **cold** and **warm**: $2/3 - 0 = 2/3$

Distance for multivalent categorical attributes

- **Jaccard Similarity**
- Instead of calculating distances between vectors, we will work with sets (unordered collection of objects).
- Jaccard similarity is not a distance, 1-JS is a distance. Can be used for multivalent categorical variables.
- Let's say we have two sets of objects: A and B.
 - Elements common to A and B (intersection): $A \cap B$
 - All elements in A and B (union): $A \cup B$.
- So the Jaccard Similarity (JS) is given by:
$$JS = |A \cap B| / |A \cup B|$$
- Example:

Movie1 (A)	Movie2 (B)	$A \cap B$	$A \cup B$	distance 1-JS
Comedy	Comedy, Action	1	2	$\frac{1}{2}$
Comedy, Action	Comedy, Action	2	2	0
Comedy, Action	Comedy, Drama	1	3	$\frac{1}{3}$
Comedy, Action	Drama, Non-fiction	0	4	1

- Euclidean distance between animals

Animal	Egg-laying	Scales	Poisonous	Slender	Legs	Reptile?	Target
	Rattlesnake	True	True	True	True	0	Yes
	Boa constrictor	False	True	False	True	0	Yes
	Dart Frog	True	False	True	False	4	No
	Alligator	True	True	False	True	4	Yes

Feature vectors for data instances (animals):

$$\text{Rattlesnake} = [1, 1, 1, 1, 0]$$

$$\text{Boa constrictor} = [0, 1, 0, 1, 0]$$

$$\text{Dart Frog} = [1, 0, 1, 0, 4]$$

$$\text{Alligator} = [1, 1, 0, 1, 4]$$

Ref: Example taken from MIT Opencourseware – 6.0002 Introduction to ML

Example I – cont'd

- Euclidean distance between animals

```
data = [ [1,1,1,1,0], [0,1,0,1,0], [1,0,1,0,4], [1,1,0,1,4] ]
cols = ['Rattlesnake', 'Boa Constrictor', 'DartFrog', 'Alligator']
df = pd.DataFrame(euclidean_distances(data, data), columns=cols, index=cols)
```

	Rattlesnake	Boa Constrictor	DartFrog	Alligator
Rattlesnake	0.000000	1.414214	4.242641	4.123106
Boa Constrictor	1.414214	0.000000	4.472136	4.123106
DartFrog	4.242641	4.472136	0.000000	1.732051
Alligator	4.123106	4.123106	1.732051	0.000000

- Alligator closer to Dart Frog than to snakes. Why?
- Alligator differs from Frog in 3 features, from Boa in only 2
- But scale on "Legs" is from 0 to 4, on other features is 0 to 1
- So "Legs" dominate as its dimension is disproportionately large

Example I – cont'd

- What if "Legs" converted to "4-Legs?" and used as a binary feature?

	Egg-laying	Scales	Poisonous	Slender	Legs	4-Legs?
Rattlesnake	1	1	1	1	0	No
Boa Constrictor	0	1	0	1	0	No
DartFrog	1	0	1	0	1	Yes
Alligator	1	1	0	1	1	Yes

	Rattlesnake	Boa Constrictor	DartFrog	Alligator
Rattlesnake	0.000000	1.414214	1.732051	1.414214
Boa Constrictor	1.414214	0.000000	2.236068	1.414214
DartFrog	1.732051	2.236068	0.000000	1.732051
Alligator	1.414214	1.414214	1.732051	0.000000

- Now Alligator is closer to snakes than it is to Dart Frog, which makes more sense.
- This shows why Feature Engineering is an important concept.

Example II

Attribute	Person1	Person2	...	PersonN	
Gender	Male	Female		Male	→ Binary
Residential status	rent	owner	...	other	→ Multi-class (3)

- Distance between Person1-Person2, Person1-PersonN, etc?
- Assign a binary dummy variable to each value of "Residential status" and compute the average distances:

$$\text{Person1} = [1, (1, 0, 0)]$$

$$\text{Person2} = [0, (0, 1, 0)]$$

$$\text{PersonN} = [1, (0, 0, 1)]$$

$$\text{Distance: Person}[1-2] = [1, 2/3] = (1+2/3)/2 = 5/6$$

$$\text{Person}[1-N] = [0, 2/3] = (0+2/3)/2 = 1/3$$

$$\text{Person}[2-N] = [1, 2/3] = (1+2/3)/2 = 5/6$$

simple average: a weight
of 1 for each feature

distance
between two
genders

$d_{i,j} = (p - m) / p$
m: # of matches
p: # of variables

- Heterogeneous attributes and differences in scale

Attribute	Person1	Person2	...	PersonN
Gender	Male	Female		Male
Age	29	47		38
Residential status	rent	owner		other
Salary	36,000	140,000		90,000

The diagram illustrates two issues with the data. Issue 1, represented by a red circle, highlights the large numerical difference in salary between Person1 (36,000) and Person2 (140,000). Issue 2, represented by a blue bracket, highlights the mixed nature of the attributes (Gender, Age, Residential status) and the large numerical difference in age between Person1 (29) and PersonN (38).

- Issues:
 - 1 Numerical variables with very different scales
 - Need to apply scaling (normalization)
 - Without scaling, our distance function will treat 10 TL difference in salary as significant as 10 yrs difference in age
 - 2 Distance issue with attributes of mixed nature
 - For distance computations, "gender" and "age" variables must be combined numerically.
So how do we do this?

Issues with kNN: Scaling

- **Data Scaling**
- As the distance between data points is essential in kNN algorithm, attribute scaling is critical.
- Min-Max scaling [0-1 range]: $x_s = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$

ID	Gender	Age	Salary
1	F	27	19,000
2	M	51	64,000
3	M	52	100,000
4	F	33	55,000
5	M	45	45,000

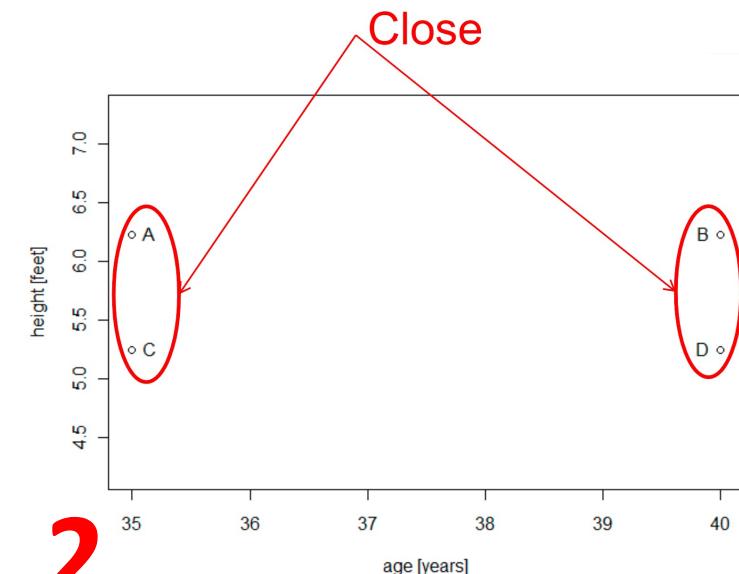
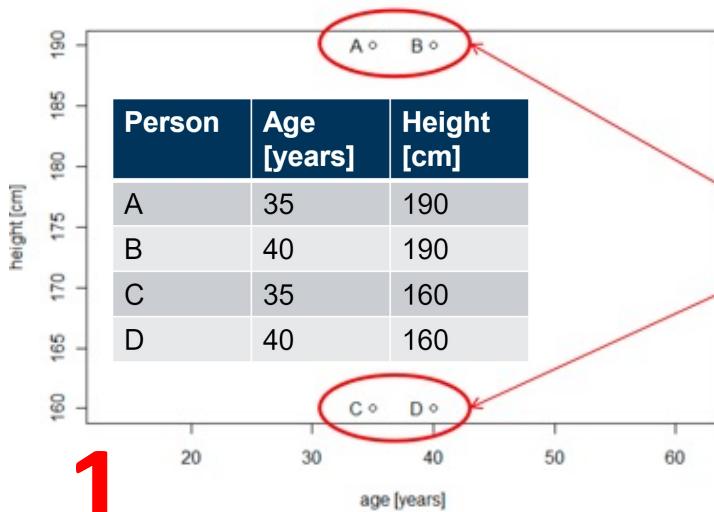
ID	Gender	Age	Salary
1	1	0.00	0.00
2	0	0.96	0.56
3	0	1.00	1.00
4	1	0.24	0.44
5	0	0.72	0.32

- Normalizing (standardization) data:

$$x_s = \frac{x - \bar{x}}{\sigma} \quad (\text{0 mean, unit variance})$$

Issues with kNN: Scaling

- Why the need for scaling?



3

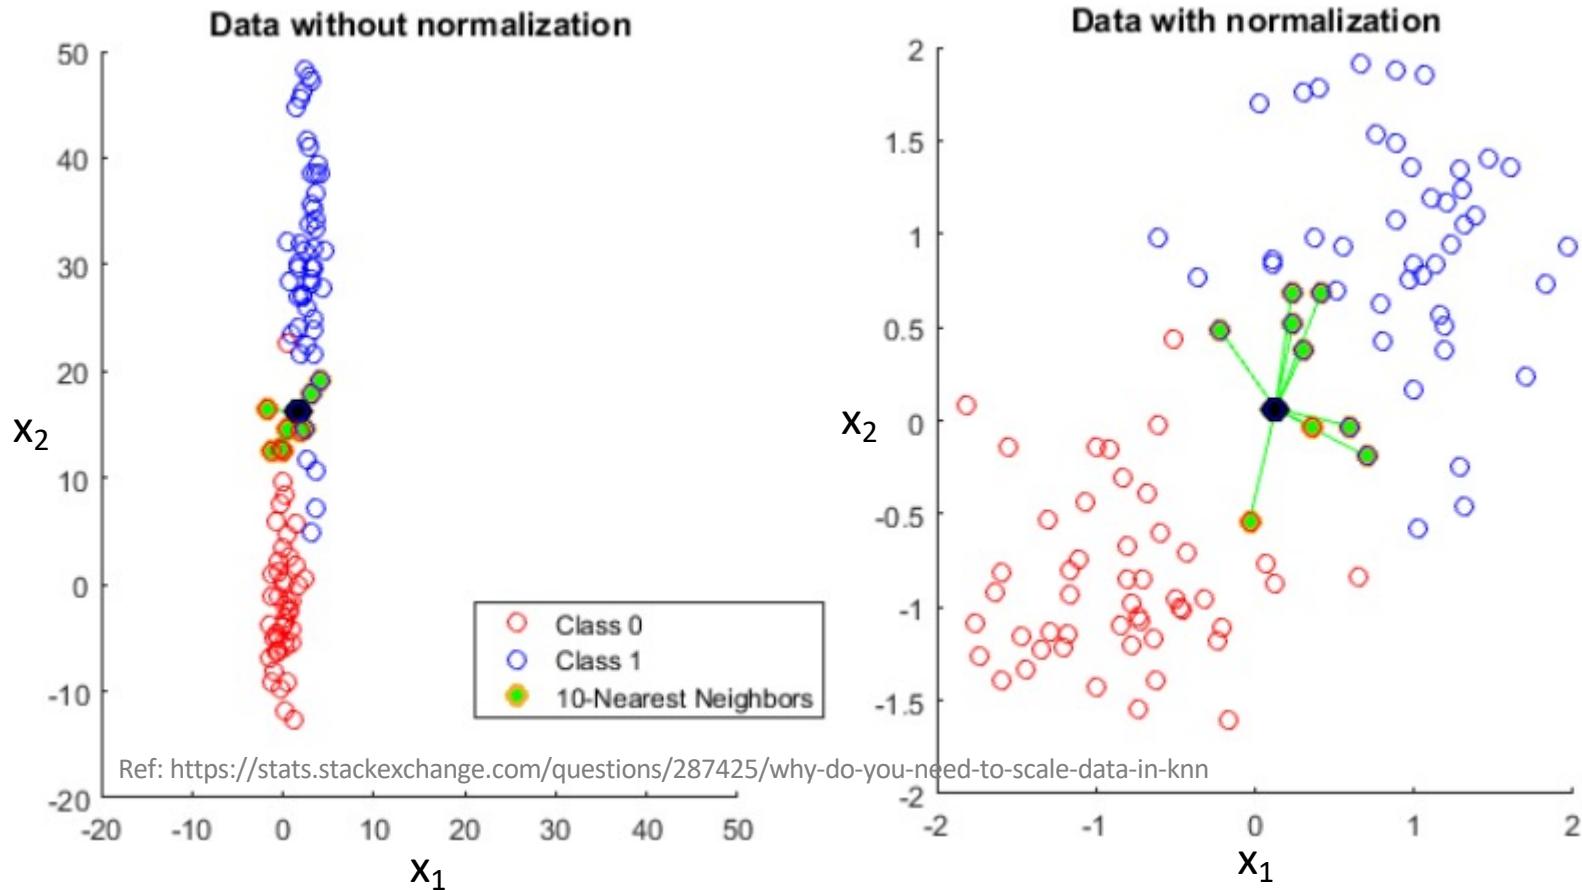
A scatter plot showing height [scaled] on the y-axis (ranging from -0.5 to 0.5) versus age [scaled] on the x-axis (ranging from -1.5 to 1.5). The same four points are plotted: A (-0.87, 0.87), B (0.87, 0.87), C (-0.87, -0.87), and D (0.87, -0.87). The points are now correctly aligned along a diagonal line, indicating that the data has been properly scaled for kNN.

Person	Age [scaled]	Height [scaled]
A	-0.87	0.87
B	0.87	0.87
C	-0.87	-0.87
D	0.87	-0.87

Ref: <https://stat.ethz.ch/education/semesters/ss2012/ams/slides/v4.2.pdf>

Issues with kNN: Scaling

- Why the need for scaling?



Consider a simple binary classification problem, where a Class 1 sample is chosen (black) along with its 10-nearest neighbors (filled green). Without normalization, all the nearest neighbors are aligned in the direction of the axis with the smaller range, i.e., x_1 leading to incorrect classification. Normalization solves this problem.

Issues with kNN: Scaling

- Why the need for scaling?

Unscaled Data

Age	Loan	Default	Distance
25	\$40,000	N	102000
35	\$60,000	N	82000
45	\$80,000	N	62000
20	\$20,000	N	122000
35	\$120,000	N	22000
52	\$18,000	N	124000
23	\$95,000	Y	47000
40	\$62,000	Y	80000
60	\$100,000	Y	42000
48	\$220,000	Y	78000
33	\$150,000	Y	80000

2

3

1

Scaled Data

Age	Loan	Default	Distance
0.125	0.11	N	0.7652
0.375	0.21	N	0.5200
0.625	0.31	N	0.3160
0	0.01	N	0.9245
0.375	0.50	N	0.3428
0.8	0.00	N	0.6220
0.075	0.38	Y	0.6669
0.5	0.22	Y	0.4437
1	0.41	Y	0.3650
0.7	1.00	Y	0.3861
0.325	0.65	Y	0.3771
0.7	0.61	?	N

Standardized Variable

$$X_s = \frac{X - \text{Min}}{\text{Max} - \text{Min}}$$

closest ~ most similar

$$D = \sqrt{(Age_1 - Age_2)^2 + (Loan_1 - Loan_2)^2}$$

Euclidean Distance

closest ~ most similar

Ref: https://www.saedsayad.com/k_nearest_neighbors.htm

Issues with kNN: Scaling

- If variables are not scaled:
 - variable with largest range has most weight
 - distance depends on scale
- Scaling gives every variable equal weight
- Similar alternative is re-weighting:

$$d(i,j) = \sqrt{w_1(x_{i1} - x_{j1})^2 + w_2(x_{i2} - x_{j2})^2 + \dots + w_p(x_{ip} - x_{jp})^2}$$

- Scale if,
 - variables measure different units (kg, meter, sec,...)
 - you explicitly want to have equal weight for each variable
- **Don't scale if units are the same for all variables**
- Most often: Better to scale (doesn't hurt)

Issues with kNN: Attributes of mixed type

- Ideally the distance between two points should be:

$$d(i,j) = \sqrt{w_1(x_{i1} - x_{j1})^2 + w_2(x_{i2} - x_{j2})^2 + \dots + w_p(x_{ip} - x_{jp})^2}$$

- It is, however, hard to decide weights on the features. Scaling takes care of this for numerical attributes, but what about in the presence of mixed attributes?
- Example: Suppose we need to predict if a person is obese or not given the height, weight and gender (Male/Female). The gender is probably very important for this problem.
- Example: For predicting the income of a person given age, education and the gender, the gender should have no bearing on the target with 0 or a very small weight.
- It's crucial to give appropriate weights (domain expertise?) to numerical and nominal attributes. See the next page for how.

Ref: <https://www.quora.com/How-do-I-perform-a-KNN-algorithm-with-a-mix-of-Categorical-and-Continuous-Variables>

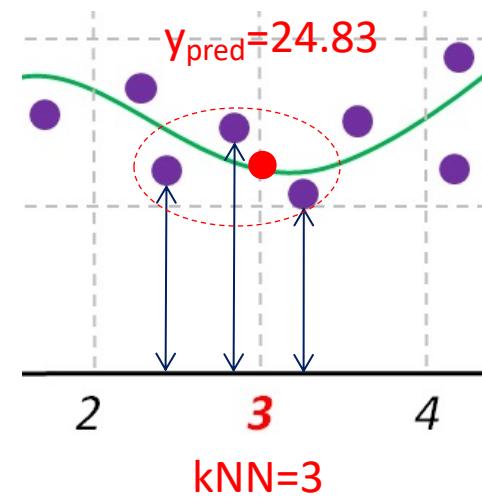
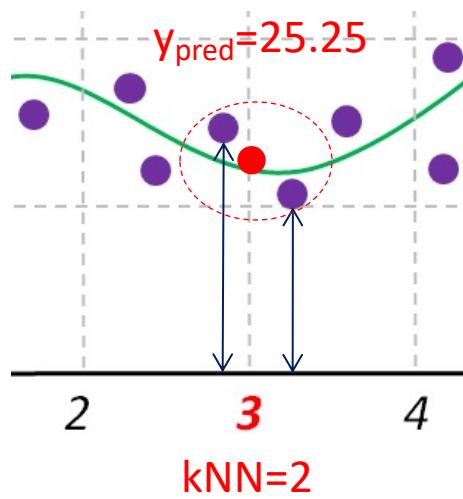
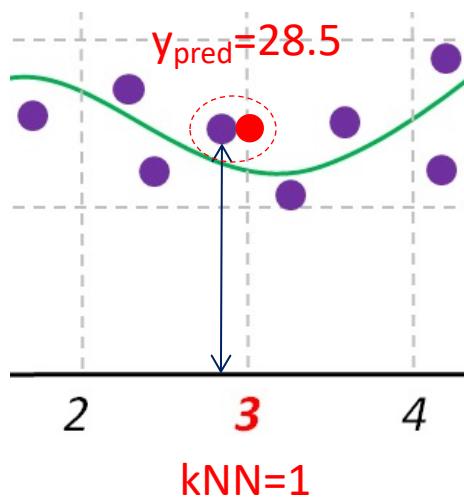
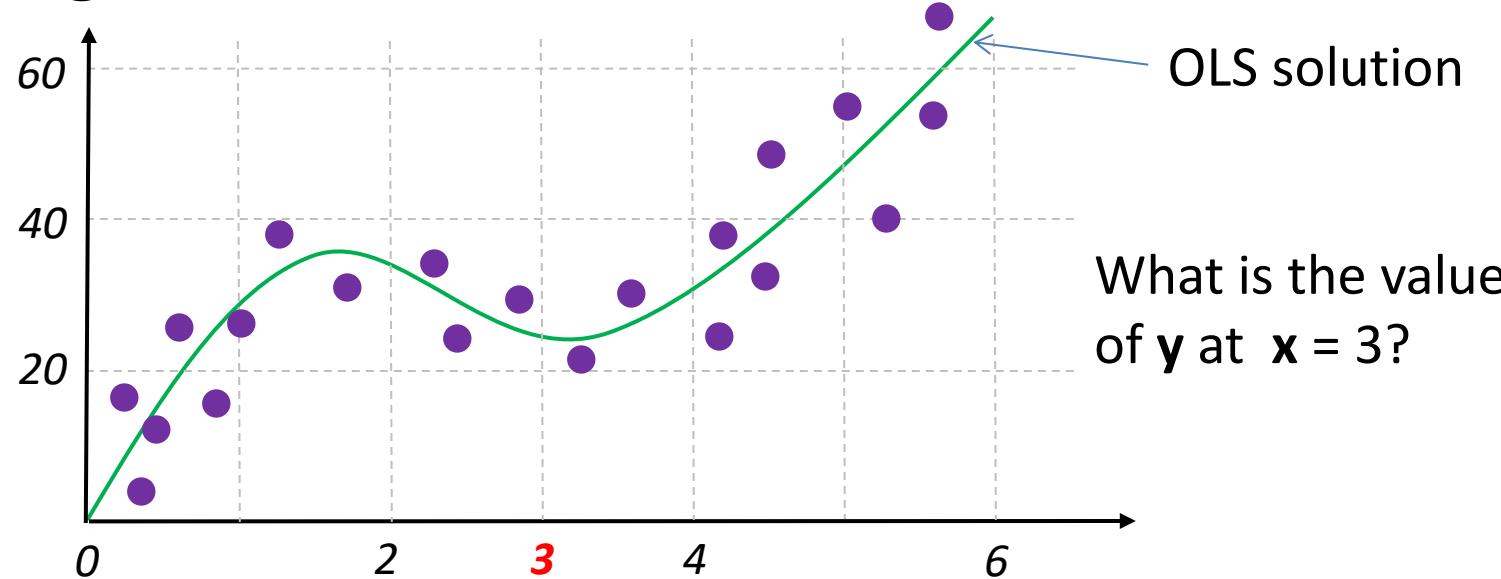
Issues with kNN: Attributes of mixed type

- Nominal (binary or otherwise), ordinal and numeric
- One popular hybrid distance function is Gower's distance (a weighted formula for combined effects):
 - where w_{ijk} is a **user-specified weight for attributes** $i,j=1,2,\dots,k$. In the absence of specific information, the weights are often set to 1 (features with equal weight).
 - There is not a generally applicable solution to choosing weights effectively (one setting that works well for one scenario may not work well for others)
 - If variable k is numeric: Use the normalized distance
 - If variable k is nominal/ordinal: Use Hamming, Simple Matching, etc.

$$d_{ij} = \frac{\sum_{k=1}^p w_{ijk} d_{ijk}}{\sum_{k=1}^p w_{ijk}}$$

Applications of kNN (1)

- Regression via kNN

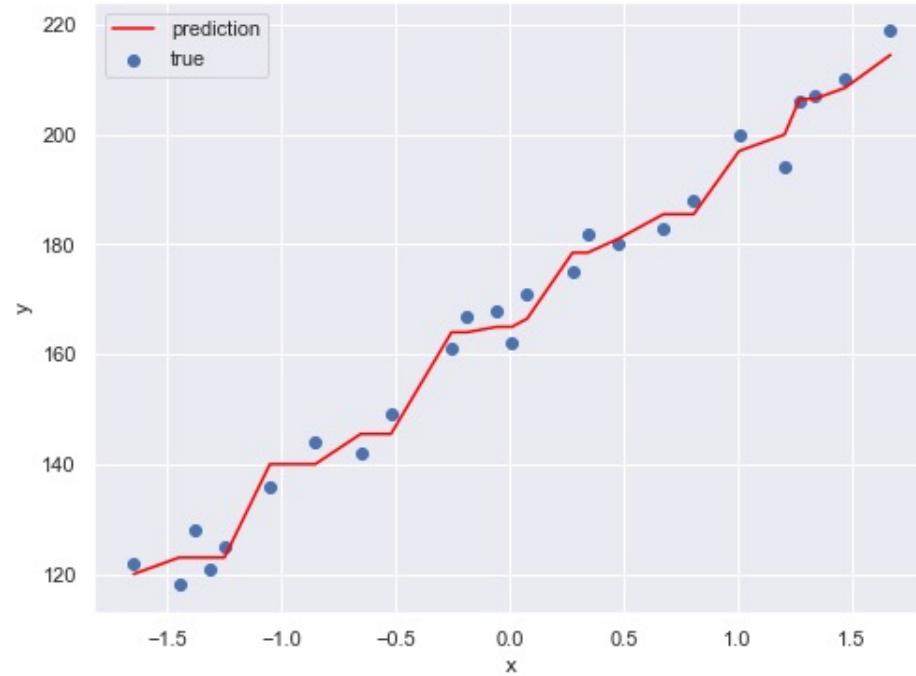


Applications of kNN (1) – cont'd

- **Regression via kNN**

```
y = [122,118,128,121,125,136,144,142,149,161,167,168,162,
      171,175,182,180,183,188,200,194,206,207,210,219],
x = [50,53,54,55,56,59,62,65,67,71,72,74,75,76,79,80,82,85,87,90,93,94,95,97,100]
```

```
kfold = RepeatedKFold(n_splits=5, n_repeats=3, random_state=42)
y = StandardScaler().fit_transformx)
knn_reg = KNeighborsRegressor()
params = {'n_neighbors': [1,2,3,4,5], 'p': [1, 2]}
grid = GridSearchCV(estimator=knn_reg, param_grid=params, cv=kfold)
grid.fit(x,y)
y_pred = grid.predict(x)
plt.scatter(x,y)
plt.plot(x, y_pred);
```



- **Missing value imputation via kNN**

- We have a data set **D** of size **N**
- Assume data in instance 'i' has a missing value for attribute **C₁**

Pseudocode for the imputation

- for $j = 1, N$ (where $j \neq i$) :
 - Compute distance d_j between the data instances 'i' and 'j'

#	C₁	C₂	C₃	C₄	Class
1	21	6.5	52	41	Yes
2	17	11	87	15	No
...
i	?	7	54	39	Yes
...
N	40	14	99	82	No

$$d(i, j) = d_j = \sqrt{\sum_{m \in D^*} (x_i^m - x_j^m)^2}$$

D^* is the set of attributes with non-missing values

where sum is over all columns with no missing values

- Save the distance d_j in a similarity array **S**
- Sort the array **S** in descending order
- Pick the top **k** data instances from **S**
- Impute the missing value of 'i' by the mean (or mode in case of categorical attribute) of the top **k** known values of attribute **C₁**

Pros and Cons of kNN

- **Pros**

- Can be used for both classification and regression
- Makes no assumptions about the data
- Simple to understand, highly intuitive, and easy to implement
- Flexible to feature/distance choices
- Naturally handles (and particularly powerful for) multiple classes
- Can do well in practice with enough representative data
- Good performance on data with few dimensions
- Easy to update in online settings (could be very slow)
- Learns nonlinear boundaries

- **Cons**

- May overfit (low bias/high variance model, choice of k is crucial)
- Missing data need to be handled
- Do not work well with high dimensional datasets
- Slow and cost of classification can be high (large search problem to find nn)
- Sensitive to outliers and irrelevant attributes, vulnerable to noisy inputs
- Need to store all samples and may require lots of memory
- No variable selection of any kind
- Using all attributes to compute distances is computationally expensive
- When data is imbalanced, predicted values are biased towards majority class
- Not a good option for mixed data (categorical and numerical) sets

```
# using scikit-learn library
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

# kcv is the optimum k value found via CV
Clf = KNeighborsClassifier(n_neighbors=kcv)
Clf.fit(X_train, y_train)
accuracy_score(y_test, Clf.predict(X_test))

# Hyperparameters tuned:
# n_neighbors: number of neighbors (def=5)
# metric: metric to use for computing distance
```